

```


import pandas as pd
import numpy as np
import re
import nltk
from nltk.corpus import stopwords
from sklearn.model_selection import train_test_split, RandomizedSearchCV, cross_val_score
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_selection import SelectKBest, chi2
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline, FeatureUnion
from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, roc_auc_score, roc_curve
import matplotlib.pyplot as plt
import seaborn as sns
from wordcloud import WordCloud
import shap

```

```

## Download resources
nltk.download('stopwords')
stop_words = set(stopwords.words('english'))

```

 [nltk\_data] Downloading package stopwords to /root/nltk\_data...  
 [nltk\_data] Unzipping corpora/stopwords.zip.

```

# 1. LOAD AND CLEAN DATA
# -----
df = pd.read_csv("/content/WELFake_Dataset.csv", on_bad_lines="skip", engine="python")
df = df.dropna(subset=['label'])
df = df[df['label'].astype(str).str.strip().isin(['0', '1'])]
df['label'] = df['label'].astype(int)
df['title'] = df['title'].fillna('').astype(str)
df['text'] = df['text'].fillna('').astype(str)
df['content'] = (df['title'] + " " + df['text']).fillna('').astype(str)

```

```

# 2. CLEAN TEXT
# -----
def clean_text(text):
    text = str(text).lower()
    text = re.sub(r"http\S+|www.\S+", '', text)
    text = re.sub(r"<.*?>", '', text)
    text = re.sub(r"^[a-z\s]", '', text)
    words = text.split()
    words = [w for w in words if w not in stop_words]
    return " ".join(words)

df['content'] = df['content'].apply(clean_text)

```

```

# 3. ADD HANDCRAFTED FEATURES
# -----
class TextStats(BaseEstimator, TransformerMixin):
    def fit(self, x, y=None):
        return self
    def transform(self, data):
        # basic stats
        return pd.DataFrame({
            'char_count': data.apply(len),
            'word_count': data.apply(lambda x: len(x.split())),
            'sentence_count': data.apply(lambda x: x.count('.')),
            'uppercase_count': data.apply(lambda x: sum(1 for c in x if c.isupper())),
            'num_count': data.apply(lambda x: sum(1 for c in x if c.isdigit())),
            'punct_count': data.apply(lambda x: sum(1 for c in x if c in "!?"))
        })

```

```
# 4. SETUP FEATURES (TF-IDF WITH N-GRAMS + HANDCRAFTED)
# -----
tfidf = TfidfVectorizer(ngram_range=(1,2), max_df=0.8, min_df=3, max_features=12000)

feature_union = FeatureUnion([
    ('tfidf', tfidf),
    ('stats', TextStats())
])

# 5. PREPROCESSING & FEATURE SELECTION PIPELINE
# -----
X_features = feature_union.fit_transform(df['content'])
# Feature selection (keep top 8000 overall for speed)
selector = SelectKBest(chi2, k=min(8000, X_features.shape[1]))
X_selected = selector.fit_transform(X_features, df['label'])
y = df['label']

# 6. SPLIT DATA
# -----
X_train, X_test, y_train, y_test = train_test_split(X_selected, y, test_size=0.2, random_state=42, stratify=y)

# 7. MODELS FOR COMPARISON
# -----
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000, random_state=42, n_jobs=-1),
    "Random Forest": RandomForestClassifier(n_estimators=100, random_state=42, n_jobs=-1),
    "Naive Bayes": MultinomialNB(),
    "Decision Tree": DecisionTreeClassifier(random_state=42),
}

results = {}
for name, clf in models.items():
    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    results[name] = acc
    print(f"\n{name}-")
    print(f"Accuracy: {acc:.4f} | ROC AUC: {roc_auc_score(y_test, y_pred):.3f}")
    print(classification_report(y_test, y_pred))

acc_df = pd.DataFrame([results], index=['Accuracy']).T
print("\nModel Performance Comparison:\n", acc_df)
```



```
-Logistic Regression-
Accuracy: 0.9263 | ROC AUC: 0.925
```

	precision	recall	f1-score	support
0	0.93	0.91	0.92	691
1	0.92	0.94	0.93	774
accuracy			0.93	1465
macro avg	0.93	0.93	0.93	1465
weighted avg	0.93	0.93	0.93	1465

```
-Random Forest-
Accuracy: 0.9420 | ROC AUC: 0.942
```

	precision	recall	f1-score	support
0	0.93	0.95	0.94	691
1	0.95	0.94	0.94	774
accuracy			0.94	1465
macro avg	0.94	0.94	0.94	1465
weighted avg	0.94	0.94	0.94	1465

```
-Naive Bayes-
Accuracy: 0.8792 | ROC AUC: 0.879
```

	precision	recall	f1-score	support
0	0.87	0.87	0.87	691

	1	0.89	0.89	0.89	774
accuracy				0.88	1465
macro avg	0.88	0.88	0.88	0.88	1465
weighted avg	0.88	0.88	0.88	0.88	1465

-Decision Tree-

Accuracy: 0.9167 | ROC AUC: 0.917

		precision	recall	f1-score	support
	0	0.91	0.92	0.91	691
	1	0.93	0.91	0.92	774

accuracy				0.92	1465
macro avg	0.92	0.92	0.92	0.92	1465
weighted avg	0.92	0.92	0.92	0.92	1465

Model Performance Comparison:

	Accuracy
Logistic Regression	0.926280
Random Forest	0.941980
Naive Bayes	0.879181
Decision Tree	0.916724

# 8. ENSEMBLE VOTING CLASSIFIER

```
# -----
voting_clf = VotingClassifier([
    ('lr', LogisticRegression(max_iter=1000, random_state=42, n_jobs=-1)),
    ('rf', RandomForestClassifier(n_estimators=100, random_state=42, n_jobs=-1)),
    ('nb', MultinomialNB())
], voting='soft', n_jobs=-1)
voting_clf.fit(X_train, y_train)
voting_pred = voting_clf.predict(X_test)
print("\n--Voting Classifier--")
print("Accuracy:", accuracy_score(y_test, voting_pred))
print("ROC AUC:", roc_auc_score(y_test, voting_pred))
print(classification_report(y_test, voting_pred))
```



```
--Voting Classifier--
Accuracy: 0.9228668941979522
ROC AUC: 0.9223469338149781
```

		precision	recall	f1-score	support
	0	0.92	0.91	0.92	691
	1	0.92	0.93	0.93	774

accuracy				0.92	1465
macro avg	0.92	0.92	0.92	0.92	1465
weighted avg	0.92	0.92	0.92	0.92	1465

# 9. CROSS-VALIDATION FOR ENSEMBLE

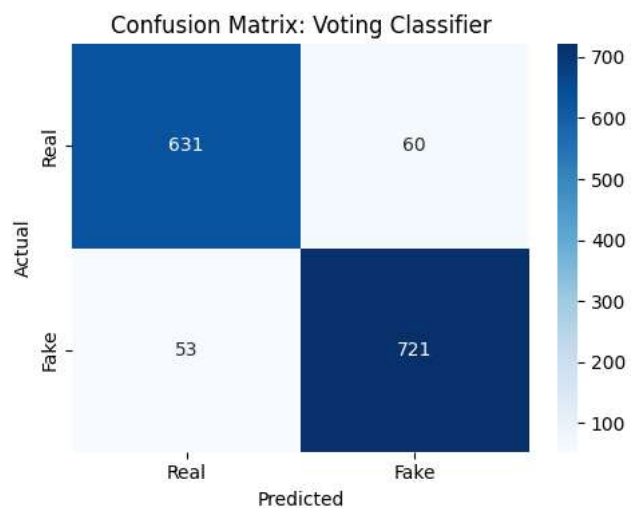
```
# -----
cv_accuracy = cross_val_score(voting_clf, X_train, y_train, scoring='accuracy', cv=3)
cv_roc = cross_val_score(voting_clf, X_train, y_train, scoring='roc_auc', cv=3)
print(f"\nVoting CV Accuracy: {cv_accuracy.mean():.3f} (+/- {cv_accuracy.std():.3f})")
print(f"Voting CV ROC AUC: {cv_roc.mean():.3f} (+/- {cv_roc.std():.3f})")
```



```
Voting CV Accuracy: 0.905 (+/- 0.005)
Voting CV ROC AUC: 0.971 (+/- 0.002)
```

# 10. CONFUSION MATRIX FOR VOTING CLASSIFIER

```
# -----
cm = confusion_matrix(y_test, voting_pred)
plt.figure(figsize=(5,4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=["Real", "Fake"], yticklabels=["Real", "Fake"])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix: Voting Classifier")
plt.tight_layout()
plt.show()
```



# 11. WORD CLOUD VISUALIZATION

# -----

```
for label, color, title in [(1, 'Reds', 'Fake News'), (0, 'Greens', 'Real News')]:
    txt = " ".join(df[df['label'] == label]['content'])
    wc = WordCloud(width=800, height=400, background_color='white', colormap=color).generate(txt)
    plt.figure(figsize=(10,5))
    plt.imshow(wc, interpolation='bilinear')
    plt.axis('off')
    plt.title(f"Most Common Words in {title}")
    plt.tight_layout()
    plt.show()
```

[illegible]


```

if hasattr(models["Random Forest"], "feature_importances_"):
    importances = models["Random Forest"].feature_importances_
    idx_top = np.argsort(importances)[-20:][::-1]
    feature_names = np.array([f"tfidf:{f}" for f in tfidf.get_feature_names_out()] + \
                             list(['char_count', 'word_count', 'sentence_count', 'uppercase_count', 'num_count', 'punct_count'])))
    top_feat_names = feature_names[selector.get_support(indices=True)][idx_top]
    plt.figure(figsize=(10,5))
    plt.barh(top_feat_names, importances[idx_top], color='teal')
    plt.xlabel("Importance")
    plt.title("Top 20 Features: Random Forest")
    plt.gca().invert_yaxis()
    plt.tight_layout()
    plt.show()

```



Top 20 Features: Random Forest



# 13. ERROR ANALYSIS: False Positives/Negatives

# -----

X\_test\_idx = y\_test.reset\_index(drop=True)

false\_pos\_idx = np.where((y\_test==0)&(voting\_pred==1))[0]

false\_neg\_idx = np.where((y\_test==1)&(voting\_pred==0))

print("\nSample False Positives (Real misclassified as Fake):")

print(df.iloc[X\_test\_idx.index[false\_pos\_idx]][['title', 'text']].head(2))

print("\nSample False Negatives (Fake misclassified as Real):")

print(df.iloc[X\_test\_idx.index[false\_neg\_idx]][['title', 'text']].head(2))